

Below is a **high-level summary** of the multi-agent solution for **decentralized drone coverage** with **truly local rewards**. We break down **what** each file does, **how** it does it, and provide some background on the **theory** underpinning this approach.

Solution Overview

We are tackling a **multi-drone coverage** task on an $N \times NN \times NN \times N$ grid. Each drone has a coverage radius, and we measure how many cells of the grid each drone covers. The **key** aspects here are:

1. **Decentralized Setting:** Each drone is controlled by **its own** DQN agent.
2. **Local Observations:** Each drone observes **only its own** position (row, col) plus a global coverage fraction (optional).
3. **Local Rewards:** Each drone's reward is how many *new* cells it covers (i.e., cells that weren't already covered by other drones).
4. **Learning Algorithm:** Each drone trains a **DQN** policy using a **neural network** for function approximation, storing transitions in a **replay buffer** and performing **Q-learning** style updates.

This setup encourages drones to spread out so that each can earn a local reward from covering unique areas of the grid.

File Summaries

1) multi_drone_env.py

- **Purpose:** Defines the multi-drone coverage environment. In a standard single-agent Gym environment, you have one observation, one action, and one reward. Here, we have a **multi-agent environment** returning:
 - A **dict** of observations: obs[i] for drone iii.
 - A **dict** of rewards: reward[i] for drone iii.
 - A **dict** of actions is expected: action[i] for drone iii.
- **Key Components:**

1. **Initialization (__init__):**
 - Sets up a $N \times NN \times NN \times N$ grid (grid_size).
 - Creates num_drones drones, each with a coverage radius (coverage_radii).
 - Sets a maximum step limit (max_steps).
 - Defines the **action space** as a dict of size num_drones, each drone having 5 discrete moves.

- Defines the **observation space** as a dict of dimension (row, col, coverage_frac) for each drone.
2. **Reset** (reset()):
- Randomly positions each drone in the grid.
 - Returns the initial **dictionary** of observations, keyed by drone ID.
3. **Step** (step(action_dict)):
- Each drone's action is interpreted and executed (moving the drone).
 - Coverage sets are computed for each drone.
 - The environment calculates **local rewards**: each drone's reward is how many new cells it uniquely covers.
 - Checks if the episode is done (full coverage or step limit).
 - Returns {drone_id: observation}, {drone_id: local_reward}, done, and info.
4. **Render** (render()):
- Prints a simple textual representation of the grid.
 - Marks coverage with "*" and drones with labels like "D0", "D1", etc.

- **Local Reward Computation Theory:**

- For each drone i , define:

$\text{coverage}_i = \{(r,c) \mid \text{dist}((r,c), \text{drone}_i) \leq \text{radius}_i\}$. $\text{coverage}_i = \{(r,c) \mid \text{dist}((r,c), \text{drone}_i) \leq \text{radius}_i\}$.

- Let $\text{coverage_minus}_i = \bigcup_{j \neq i} \text{coverage}_j$. $\text{coverage_minus}_i = \bigcup_{j \neq i} \text{coverage}_j$.

- **Newly contributed cells** by drone i is

$\text{new}_i = \text{coverage}_i \setminus \text{coverage_minus}_i$. $\text{new}_i = \text{coverage}_i \setminus \text{coverage_minus}_i$.

- The reward for drone i is $|\text{new}_i|$ (possibly normalized by the total grid size). This ensures each drone only gets credit for coverage that *wouldn't* exist without it.

2) agent.py

- **Purpose:** Implements a **DQN agent** that each drone uses to decide actions and learn via replay-buffer Q-learning.
- **Key Components:**
 1. **Initialization** (__init__):

- Builds a neural network (policy_net) mapping **observations** (3(3(3-dim)→) \to)→ **Q-values** (5(5(5-dim())
- Creates an **Adam optimizer**, a **mean squared error** loss, and sets up a **replay buffer**.

2. Action Selection (act(obs)):

- ϵ -greedy: pick a random action with probability ϵ , else pick the action that maximizes the Q-network.

3. Memory Storage (step(obs, action, reward, next_obs, done)):

- Stores transitions in a **replay buffer** for offline training.

4. Learning (learn()):

- Samples **minibatches** from the replay buffer.
- **Current Q**: $Q_{\theta}(\text{obs}, \text{action})$
- **Target Q**: $\text{reward} + \gamma \max_{a'} Q_{\theta}(\text{next_obs}, a')$ if not done.
- Minimizes MSE loss between current and target Q via **backprop**.
- Decays ϵ after each training step.

• DQN Theory:

- **Q-Learning** is a temporal-difference method:

$Q(s,a) \leftarrow Q(s,a) + \alpha [r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$. $Q(s,a) \rightarrow Q(s,a) + \alpha [r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$.

- **DQN** (Deep Q-Network) parameterizes Q_{θ} with a neural net.
- **Replay Buffer** ensures i.i.d. training data and stabilizes learning.
- This agent code is replicated for **each** drone, so each drone has **its own** DQN.

3) training.py

- **Purpose:** The **coordinator** script that instantiates the environment and the multiple DQN agents (one per drone). It then runs a training loop over multiple episodes.
- **Key Components:**
 1. **train_multi_drone(...)**
 - Creates the environment (MultiDroneCoverageEnv).
 - Creates one DroneDQNAgent per drone.
 - For each **episode**:

- **Reset** the env, get initial obs_dict.
 - For each time-step until done:
 1. Ask each agent for an action based on its local observation.
 2. Call env.step(action_dict), receiving (next_obs_dict, reward_dict, done, info).
 3. Each agent stores (obs, action, reward, next_obs, done) in its own replay buffer.
 4. Each agent calls learn() once, sampling from its buffer.
 5. Update obs_dict = next_obs_dict.
 - Track average local reward (or some metric) and print progress.
 - Returns the trained agents and the reward history.
2. **demo_run(...)**
- Runs a single test episode with the learned policies (still ϵ -greedy, but you could set $\epsilon=0$ for pure exploitation).
 - Prints out the environment's textual rendering each step.

- **Multi-Agent Reinforcement Learning Theory:**

- We're using a **decentralized** approach: each agent has its **own** policy and sees **its own** observation.
- Agents have separate replay buffers and do not share parameters or experiences. They learn in **parallel** within the same environment.
- The environment dispatches local rewards, so each agent's objective is to maximize its own discounted return. In this scenario, local reward is newly contributed coverage.
- In principle, the combination of local rewards can also lead to a good global coverage if the drones learn to coordinate by not overlapping too much (since overlapping coverage yields zero additional local reward).

Technical/ Theoretical Details

1. State Space & Observations

- Each drone's observation is a 3D vector: (row,col,coverage_frac)(row, col, coverage_frac)(row,col,coverage_frac).
- If we wanted a more **realistic** partial-map input, we could expand the observation with local coverage data or sensor readings.

2. Action Space

- Each drone chooses one of 5 moves: UP, DOWN, LEFT, RIGHT, or STAY.
- The environment combines these into a **dict** of actions to step all drones simultaneously.

3. Local Rewards for Coverage

- The environment determines for each drone i the set of cells it covers (coverage _{i}), then subtracts the union of coverage from all other drones to find *newly contributed coverage*.
- This local reward structure is *shaping* each drone to **maximize** its unique coverage area.

4. DQN / Q-Learning

- Each drone's policy π_i is derived from a neural net $Q_{\theta_i}(s,a)$.
- The agent picks $a = \arg\max_a Q_{\theta_i}(s,a)$ with probability $(1-\epsilon)$, or random otherwise.
- The **TD target** for each transition is $r_i + \gamma \max_{a'} Q_{\theta_i}(s',a') - Q_{\theta_i}(s,a)$.
- Over time, ϵ decays, moving from exploration to exploitation.

5. Decentralized vs. Centralized

- Here, we call it “decentralized” because each drone sees only its own local state and gets its own local reward.
- The environment is *technically centralized* in the sense that it calculates coverage sets for all drones. But from the agent's perspective, no drone sees other drones' states or shares parameters.

6. Scalability

- Because we use **function approximation** (PyTorch networks), we can, in principle, handle bigger grids than a tabular Q-table.
- However, if the grid becomes very large (e.g., 50×50, 100×100, etc.), we may need more advanced neural architectures (e.g., CNNs) or more sophisticated multi-agent RL methods (QMIX, MADDPG, etc.).

In Summary

- **multi_drone_env.py**: Multi-agent environment logic, local coverage sets, local reward calculation.
- **agent.py**: Single-agent DQN class, including neural net, replay buffer, ϵ -greedy strategy, and Q-learning updates.

- **training.py**: A coordinator script that (1) instantiates multiple DQN agents (one per drone), (2) runs training episodes (collecting transitions, doing learning steps), and (3) demonstrates the learned policies.

The overall approach is a **decentralized multi-agent DQN** solution that incentivizes drones to discover **non-overlapping coverage** because **each** drone only gains reward from the coverage it uniquely provides.